


Private Banking Discovery Questions

The bank's problem is a clock. Data lands at nine; the relationship manager needs her day ready at six: the brief, the calls, the client stories, the ideas and the tasks. These are the questions to ask so the workbench we built can be fed early enough, from the right systems, with the right entitlements, and so the agents that write content and reach clients do so within the bank's rules. Every question says why we ask it.

How we are discovering

We start from what a relationship manager needs, not from systems. Six categories of need, taken from the workbench as built, are the spine of the whole page. Section 2 lists the information, insights and tasks under each category for the client to confirm. Sections 3 and 4 ask how the RM works and who may see what. Sections 5 and 6 ask what each need is made of. Sections 7 to 9 ask where it comes from and how fresh, complete and event-driven each source is, which is where the six o'clock answer lives. Sections 10 and 11 ask what the machine may write, send and do. Section 13 shows the data model that ties it together. Every question carries the category it serves, so the client can read the page by category or by section.



C1 Client Intelligence & 360 View

Client 360 and portfolio deep dive; behavioural and factual client vector; context strip and switcher; client journey.

Latency impact high: stale positions and behavioural signals mean the 360 the RM opens at nine is already hours old.

C2 Market Signals & Risk Analytics

Signal and impact analysis; risk rubric (capacity, appetite, horizon); machine-checked confidence.

Latency impact high: signals and confidence computed on yesterday's close understate intraday risk.

C3 Ranked Actions & Trade Ideas

Combined risk, ranked actions, trade ideas; the book cockpit with Now, 7-day and 30-day lanes.

Latency impact high: a "ranked now" list is only as good as its inputs; late data reorders or omits the real priorities.

C4 Workflow, Outreach & Governance

Approvals, outreach and audit; promise ledger and idea desk; in-call guardrails and the communication gateway.

Latency impact medium: workflow is not latency-sensitive, but the content it carries can be built on stale facts.

C5 Daily Planning & Conversations

Today: morning brief, call sheet, lanes; the call plan; the assistant drawer and companion.

Latency impact high: this is the six-versus-nine problem in its purest form, the brief the RM reads before her first call.

C6 Team & Data Platform Foundation

Team page, gateway, playbooks, shadow mode; RM model, identity, row-level security; incremental and Kafka ingestion.

Foundational: the enabling layer every other category rests on.

01

The problem: 9 a.m. data, 6 a.m. need

Today the RM starts working when the data arrives. The goal is the reverse: when she opens the workbench at six, the machine has already read everything, ranked the book, written each client's story, run the signals against every household, found which clients an idea fits, drafted the messages, and lined up the tasks. The RM's morning is spent deciding, not assembling.

Three hours have to come from somewhere. Either the sources can deliver earlier or incrementally, or the workbench works from the previous close plus overnight events and refreshes when the day's batch lands. Which of those is possible is a question about each source system, and most of section 7 exists to answer it.

Today

09:00 batch lands

Positions, prices and balances arrive together; the RM reads, reconciles and plans until mid-morning.

Target

04:00 to 05:30 machine runs

Overnight events applied, vectors and alerts rebuilt, briefs and drafts written, call sheet packed, tasks proposed.

Target

06:00 RM opens Today

Brief, call sheet, client journeys, ideas and drafts are ready. Anything that changes after the day's batch is marked as refreshed.

Rule

Nothing invented to fill the gap

Where a source cannot deliver by six, the screen shows the last known value with its timestamp, never a guess.

02

What we built: information, insights and tasks to confirm

Everything the workbench produces today, grouped under the six categories of RM need and, within each, the sub-areas an RM thinks in. This is the list to put in front of the client first. For each row they mark keep, change, not needed or missing; the remaining sections then connect to the rows that survive, because every data, source, recency and event question below exists to feed one of these.

Information is a fact about the client or the market, read from the systems. Insight is a conclusion the machine draws and can explain, with a level 2 "why" and level 3 rows behind it. A task is a call to action with an owner, a due date and a defined proof of completion, proposed by the machine and confirmed by the RM.

CATEGORY	INFORMATION	INSIGHTS	TASKS AND CALLS TO ACTION	WHERE
C1 Client Intelligence & 360 View Latency impact: high				
1 • Client and household	Profile, segment, residence and tax domicile, booking centre, reporting language, KYC status and review date, PEP; portfolios and service models; last contact date and channel	Where the client stands in one paragraph, with every figure sourced; moments ahead (dated needs, KYC, maturities); sensitivities quoted from notes with dates	Schedule the KYC refresh; update the profile after a review; correct a stale record	Journey 1; client compa guardr
2 • Portfolio and performance	Household value and allocation by asset class at each snapshot; holdings with cost, weight, valuation date, liquidity tier; transactions and income	Change since baseline split into price, FX and flows; top movers; income yield; stale valuations flagged	Prepare the performance story for the call; produce the "since we last spoke" pack in the client's language	Journey 2; port rooms
C2 Market Signals & Risk Analytics Latency impact: high				
3 • Mandate and concentration	Mandate bands, targets, single-position limits, exclusion lists; direct and look-through exposure by name, sector, region, currency	Band breaches with drift versus client-directed distinction; concentration and look-through breaches with the saleable part; exclusion holdings in exclusion-bound mandates	Trim to the limit or record the client's instruction to hold; rebalance toward the bands; replace an excluded holding	Alerts; and co board; ideas

CATEGORY	INFORMATION	INSIGHTS	TASKS AND CALLS TO ACTION	WHERE
4 • Liquidity and lending	Cash needs with dates and certainty; commitments and uncalled capital; facilities with drawn, limit, lending value, LTV per snapshot and trigger levels	Twelve-month coverage of needs by daily-liquid assets; LTV headroom trend and projected breach date; margin-call proximity; scenario effect on LTV and coverage	Agree the funding route before the need; discuss collateral top-up, reduced drawing or de-risking, with a second RM check; raise cash ahead of a dated liability	Cash-f board; call pla term
5 • Risk and suitability	Stated risk profile, score and horizon; behavioural features from transactions (32 in the customer vector); peer group	Rubric: capacity, appetite, horizon scored by rules, a statistical model and a language model with a combined confidence; stated-versus-observed mismatches; combined risk cell and composite score	Re-assess after a review; override a dimension with a reason; lock the rubric as the suitability record	Rubric risk an journe 3
6 • Market signals and impact	Market event log with date, severity, transmission; market series; named stress scenarios	Which signals reach the household and through which holdings, with exposure and confidence; base and severe stress on value, LTV and coverage; the one scenario most relevant to this client	Call the most exposed clients first; discuss the impact; send a market update	Signal: impact journe 3; mor
C3 Ranked Actions & Trade Ideas Latency impact: high				

CATEGORY	INFORMATION	INSIGHTS	TASKS AND CALLS TO ACTION	WHERE
8 · Ideas and opportunities	Approved product universe; house view and research; idea rules	Trade ideas per client with rationale, suitability gate and confidence, blocked ones shown with the failed rule; which clients an idea or a signal fits across the book; opportunities: lending, discretionary mandate, succession, next generation, deployment, each with its rule or quoted note	Approve or reject an idea; draft one message per client in their language; open the lending or mandate conversation	Trade room;
C4 Workflow, Outreach & Governance Latency impact: medium				
7 · Relationship and follow-through	Notes, calls and meetings with channel and date; client messages awaiting reply; promises made by RM or client with the quoted sentence	Contact cadence against the segment policy and relationship debt; open promises and those past due; decision debt (agreed but not executed); sensitivities to avoid raising	Reply to the waiting message today; fulfil each promise with its draft attached; bring the unexecuted decision back; call the client who is overdue for contact	Promises call plan relationship term; compensation

CATEGORY	INFORMATION	INSIGHTS	TASKS AND CALLS TO ACTION	WHERE
10 · Actions, approvals and record	Ranked actions with evidence, benefit and trade-off; alerts awaiting triage; approvals awaiting a checker; the nine workflow steps	Suitability, compliance, explainability and maker-checker checks per action; what is still open for this client	Triage an alert; approve or reject an action; give or obtain the second-RM check; log the call with note, cues and promises; every decision as an audit event	Workfl journe 4 and assista propos
11 · Client communication	Reporting language and preferred channel; approved templates and disclosures; cross-border rules by residence	Draft outreach from approved actions and signals with facts used and caveats; the gateway's verdict on every message (cross-border, language, promises, suitability basis, disclosures)	Edit and send through the gateway; log the send; a "since we last spoke" summary; personalised content from the house view	Workfl outrea comm gatewa desk d
C5 Daily Planning & Conversations Latency impact: high				
9 · Priorities and the day	Open items per client; deadlines; RM capacity and hours; client time zones and channels	Items in Now, week and month lanes with momentum against the previous snapshot; client urgency; whom to call, by when, in which slot and channel, with five explained priority terms; the four-sentence morning brief	Make the calls in order; defer with a reason; mark done; schedule reviews separately from calls	Today; sheet;

CATEGORY	INFORMATION	INSIGHTS	TASKS AND CALLS TO ACTION	WHERE
12 · Conversation and assistance	The client in context on every screen; the question the RM asks in her words	Answers from the engines about a client, a group of clients or the market, with the tool trace; cue answers during a call with guardrails; playbook steps with what an agent would have said	Tasks proposed from language and confirmed with one click; navigate; add a cue to the note	Assist; draw; compare; playback
C6 Team & Data Platform Foundation Foundational				
13 · Team and oversight	RM assignments and teams; overrides with reasons; approvals waiting; coverage against cadence; call load and deferrals; fee proxy	Risk, conduct, coverage and capacity by RM; ten worst households; ideas landing (reached, drafted, sent); explainability sample	Sign what waits for a checker; coach on conduct patterns; reassign or cover; pause all agents	Team
14 · Data and machine trust	Load runs, data-quality issues, dead letters; model calls and errors; shadow grades per playbook	Data coverage and freshness per source; agreement between agent drafts and RM words, against the threshold that earns the next autonomy level	Fix or own a data issue; grade a shadow draft; raise or lower the autonomy level	Audit & quality machine

C-1 **all six** For each category: is this what your RMs need, is anything here noise, and what is missing that they ask for today?

Why we ask The rows that survive define the scope; every question in the later sections traces back to one of them.

C-2 **all six** Which insights must be ready at six, which may arrive with the nine o'clock batch and refresh, and which are only needed on demand?

Why we ask Sets the overnight run and tells us which sources in section 8 must deliver early.

C-3 **all six** Which tasks are the RM's alone, which may a machine complete, and which belong to someone else (assistant, operations, compliance)?

Why we ask Draws the autonomy line per category before any agent is built.

C-4 **all six** Which categories may be shown to the client, in what form, and which stay internal?

Why we ask Separates RM-facing insight from client-facing content and decides where the communication gateway must sit.

What the RM does

The workbench is organised around six moments of the RM's day: wake, prepare, converse, follow through, grow, reflect, with the team head's Monday above them. These questions confirm the bank's RM day matches, and what she does that we have not seen.

J-1 **Daily Planning & Conversations** Walk us through an RM's day from arrival to leaving: what she reads first, how she decides whom to call, what she prepares before a call, what she does during and after it, and what she hands to others.

Why we ask The six moments we built are a hypothesis; the bank's real sequence decides what must be ready at six and in what order.

J-2 **Daily Planning & Conversations** How many clients does an RM cover, how many conversations a day is normal, and what share of the day goes to preparation, documentation and chasing versus talking to clients?

Why we ask The call plan's capacity and the efficiency case both rest on these numbers; the workbench should remove the preparation and documentation share.

J-3 **Client Intelligence & 360** What does the RM assemble by hand today before a client call, and from how many screens or systems?

Why we ask That assembly is the client journey's five chapters; whatever she gathers manually is what the machine must pre-build.

J-4 **Daily Planning & Conversations** How does the RM decide whom to call first today, and whom she can leave until next week? Is it gut, a report, or a manager's list?

Why we ask The call plan ranks by harm, deadline, bundling and relationship debt; the bank's actual heuristics should become its policy file, not be replaced by ours.

J-5 **Workflow, Outreach & Governance** What must the RM record after every interaction (call note, suitability confirmation, disclosures given, promises made), where, and how long does it take?

Why we ask The companion drafts the note, extracts promises and pre-fills the suitability record; the bank's evidence requirements define what "complete" means.

J-6 **Workflow, Outreach & Governance** Which decisions need a second person (checker, team head, compliance), and how long do those approvals take today?

Why we ask Maker-checker is configurable in the workflow; approval latency is a task the machine can queue and chase before six.

J-7

Market Signals & Risk

What does the RM do when a market event happens intraday: who tells her, how does she find which clients are affected, and what is she expected to do by when?

Why we ask Signals mapped to holdings and the impact engine automate exactly this; the expected response time decides whether intraday events must flow, not just the overnight batch.

J-8

Team & Data Platform

What does the team head look at on Monday morning, and what does she ask RMs for by email today?

Why we ask The team page has five panels built on a guess of those questions; the real ones replace them.

J-9

Team & Data Platform

Which parts of the day would the RM most like to stop doing, and which would she never hand to a machine?

Why we ask Sets the autonomy ladder honestly: drafts only, approve to send, or informational messages sent by an agent.

Entitlements for the RM

One RM must not see another's clients. The workbench enforces this in the database with row-level security driven by dated RM-to-client assignments. These questions establish where that truth lives in the bank.

E-1

Team & Data Platform

Where is the RM-to-client assignment held (CRM, core banking, HR), and is history kept with valid-from and valid-to dates?

Why we ask Entitlement is computed from assignments valid today; the audit trail needs history to answer who could see a client in a given month.

E-2

Team & Data Platform

Besides the primary RM, who else legitimately sees a client: secondary RM, assistant, checker, team head, holiday cover, product specialist, compliance? Is each a role with a duration?

Why we ask Our assignment model has primary, secondary, checker and cover roles with dates; the bank's list decides which roles exist and whether cover is time-boxed.

E-3

Team & Data Platform

How is a team or desk defined, who heads it, and may a head see every client in the team or only aggregates?

Why we ask Team scope in row-level security is the head's whole team; if the bank wants aggregates only, the team page must change.

E-4

Team & Data Platform

Which identity provider issues RM and manager identities, and which claims carry role, team and RM id?

Why we ask The API verifies a token and maps its claims to scope; production needs the bank's issuer, audience and claim names.

E-5

Team & Data Platform

When a client is reassigned, how quickly must the old RM lose access and the new RM gain it, and who approves the reassignment?

Why we ask An assignment message updates access the moment it is applied; the bank's expectation decides whether that must be an event or can wait for a nightly file.

E-6

Team & Data Platform

May any book-wide computation happen (percentiles, peer groups, which clients fit an idea) as long as an RM only ever sees her own clients in the result?

Why we ask The analytics service computes peers and percentiles across the bank under a service role; the idea desk searches the whole book then filters to the caller's scope.

E-7

Team & Data Platform

What may an RM see about a client after offboarding or reassignment: nothing, read-only history, or her own notes only?

Why we ask Closed assignments hide the client today; the bank may want read-only access to the RM's own notes for a period.

Asset classes supported

Per asset class: the products, the client-relevant lifecycle events, and the controls in front of a transaction. The workbench classifies holdings as cash and equivalents, fixed income, equity, alternatives, commodities and structured products, with credit facilities alongside.

Cash, deposits and FX

A-C1 **Client Intelligence & 360** Which cash products do clients hold: current accounts, term deposits, structured deposits, dual-currency investments, FX accumulators? Which does the bank classify as cash and which as structured products?

Why we ask Classification drives liquidity tiers and the coverage of cash needs; a structured deposit counted as cash overstates what is sellable.

A-C2 **Client Intelligence & 360** How is available liquidity defined: cleared, uncleared, pledged, under notice, reserved for settlement? Which may the RM quote as drawable?

Why we ask The liquidity shortfall alert compares daily-liquid assets with twelve-month needs; the definition must match what the bank would actually release.

A-C3 **Client Intelligence & 360** Which FX transactions are client-initiated and which are funding conversions the bank makes? Are rate and spread visible per transaction?

Why we ask Attribution separates price, FX and flows; a bank conversion booked as a client trade misstates client behaviour in the vector.

A-C4 **Client Intelligence & 360** For deposits and accumulators: where are maturity, rollover instruction, fixings, knock-in and knock-out recorded, and when is the client told?

Why we ask Maturities and fixings are dated events the call plan should surface before they happen, not after.

Equities, ETFs and IPO

A-E1 **Client Intelligence & 360** Which service model applies per portfolio: discretionary, advisory, execution-only, custody? Who initiates orders in each and which fall under the RM's suitability duty?

Why we ask Mandate bands are measured only on managed portfolios; the flag decides which breaches are the RM's problem.

A-E2 **Market Signals & Risk** What single-position and issuer concentration limits apply at portfolio and household level, and do they look through ETFs, funds and notes to the issuer?

Why we ask Look-through concentration is the most visible alert in the workbench; the bank's limit definition must match or the alert is noise.

A-E3 **Client Intelligence & 360** How are IPO subscriptions and financing represented before listing?

Why we ask Pre-listing exposure and the loan behind it otherwise vanish from the household view and from LTV.

A-E4 **Daily Planning & Conversations** Which corporate actions reach the RM as decisions (rights, takeovers, scrip elections) versus processed silently, and by what deadline must the client elect?

Why we ask Dated elections belong in the call plan's clock term; the RM should see them on the sheet before the deadline.

A-E5 **Client Intelligence & 360** Who computes realised and unrealised profit and loss, with which cost basis and currency, and does the client statement use the same figures?

Why we ask The journey's performance story must agree with what the client holds in hand.

Funds and alternatives

A-F1 **Market Signals & Risk** Which fund types are held: daily unit trusts, money-market funds, gated or locked funds, private equity and credit with capital calls, hedge funds with monthly or quarterly dealing?

Why we ask Each maps to a liquidity tier; the coverage ratio and the stress confidence band depend on it.

A-F2 **Market Signals & Risk** For committed capital: where are commitment, called to date, uncalled and the expected call window, and how does a call notice arrive?

Why we ask Uncalled commitments are dated cash needs the RM has to fund; a call notice is a Now item.

A-F3 **Client Intelligence & 360** How stale may a fund valuation be before it is flagged, and is the valuation date carried on the position?

Why we ask The stale-valuation alert and the honesty of household totals depend on the date travelling with the number.

Bonds and structured products

A-B1 **Market Signals & Risk** Are structured notes and equity-linked notes decomposed anywhere into underlyings, barriers, coupons and autocall dates, or held as a single line with a price?

Why we ask Look-through and the impact engine need the legs; today they come from a reference file the bank should replace with its structuring system.

A-B2 **Market Signals & Risk** Which bond events reach the RM: coupon, call, maturity, rating change, restructuring? Where is the call schedule held?

Why we ask Calls and maturities are cash events; downgrades are signals against the household.

A-B3 **Market Signals & Risk** Are pledged positions marked with lending value and advance rate per line?

Why we ask Line-level lending values let the impact engine recompute LTV under a scenario and explain a margin call.

A-B4 **Ranked Actions & Ideas** Which product risk rating scheme applies and how does it map to the client's profile? Where is the appropriateness assessment for complex products recorded?

Why we ask The suitability gate on every idea and action must run the bank's rules, not ours.

Lending, collateral and mandates

A-L1 **Market Signals & Risk** Which facilities exist against a household and is each linked to its collateral portfolio?

Why we ask The facility-to-portfolio link is what makes LTV explainable line by line and lets the RM say which sale relieves it.

A-L2 **Market Signals & Risk** How often is LTV recomputed, what are trigger, top-up and close-out levels, how much notice does the client get, and who may waive?

Why we ask Headroom trend, the projected breach date and the waiver flag on the collateral board come from these.

A-L3 **Market Signals & Risk** Where are mandate bands, target weights, single-position limits and exclusion lists held, and how are client-directed deviations recorded?

Why we ask Mandate breaches feed the risk matrix; client-directed ones are a conduct item, not an RM failure, and must be distinguishable.

A-L4 **Market Signals & Risk** What is the bank's risk-profile scale, horizon questionnaire and review cycle, and is comparing stated with observed behaviour acceptable to compliance?

Why we ask The rubric scores capacity, appetite and horizon and flags stated-versus-observed mismatches; the bank must want that comparison.

Cross-asset controls

A-X1 **Workflow, Outreach & Governance** Which cross-border rules apply by client residence and booking centre: what may be discussed, offered and sent, and what needs documented reverse enquiry?

Why we ask The companion's guardrails and the communication gateway enforce a reference file; compliance owns the real matrix.

A-X2 **Workflow, Outreach & Governance** How are complaints, vulnerable-client flags, deceased or dormant statuses represented, and must they stop outreach?

Why we ask Escalation triggers and the gateway need structured flags; free text in notes is not a control.

Data, information and insights required

Three layers. Data is what the sources hold. Information is what the RM reads. Insight is what the machine concludes and must be able to explain. The workbench keeps all three visible: level 1 the conclusion, level 2 the why, level 3 the rows.

LAYER	WHAT THE RM NEEDS BY 6 A.M.	BUILT FROM	QUESTION FOR THE BANK
Data	Positions and balances per household at the last close, prices with their dates, transactions since the previous day, facility LTVs, client profile and KYC dates, notes and messages, RM assignments	Sources in section 6	Which of these can be complete before 6 a.m., which arrive at 9 a.m., and which are only available T+1? (Answered per system in section 7.)
Information	Household total and change since baseline, allocation against bands, concentration, liquidity coverage, LTV headroom, upcoming dated events, last contact and open promises	Attribution, mandate, exposure, cash-flow and facility engines	Which of these does the bank already produce (statements, reports), so the workbench should read rather than recompute? Where they disagree, which is authoritative for the client?
Insight	Ranked alerts in Now, week and month lanes; risk cell and score; which signals reached the household and by how much; rubric mismatches; ranked actions and ideas with suitability; whom to call first and why; the client's story in prose	Horizon, combined risk, signal and impact, rubric, actions, call plan, journey prose	Which insights may be shown to the RM without a human check, which need one, and which may be shown to the client? What confidence or evidence must accompany each?

I-1 **Client Intelligence & 360** Which reports does the RM receive today at nine, from which system, and what does she do with each?

Why we ask Each report is either replaced by a screen, read as an input, or retired; knowing the list stops us rebuilding what exists.

I-2 **Client Intelligence & 360** What is the baseline for "what changed": year start, last statement, last meeting, previous day?

Why we ask Attribution and momentum are computed against a baseline snapshot; the RM's expectation sets the default.

I-3 **Market Signals & Risk** Which market events, research and house views does the bank publish, in what form, and who is allowed to turn them into client-level signals?

Why we ask The signal engine maps events to holdings through transmission rules; the bank's research desk should own the event list and the rules.

I-4 **Market Signals & Risk** Which named stress scenarios does risk management publish, and may they be shown to clients?

Why we ask The journey shows one scenario per client chosen by relevance; it must be a scenario the bank stands behind.

I-5 **Ranked Actions & Ideas** What evidence must accompany a recommendation before an RM may act on it or repeat it to a client?

Why we ask Every action carries evidence, sources, suitability and confidence; the bank's standard defines the minimum.

I-6 **Team & Data Platform** Which insights would the bank want computed across the whole book for the team head, and at what granularity may they be shown?

Why we ask The team page aggregates by RM under the head's scope; conduct indicators such as override rate need thresholds the bank sets.

Source systems

One system of record per domain, one identifier that joins to the client, and a named owner. Fill the table with the bank, then ask the questions beneath for every row.

DOMAIN	TYPICAL SYSTEM	KEY	WORKBENCH USE	ASK
Client and household	CRM, client master	Client id, household id	Client record, journey chapter 1, cross-border, language	Is the household a keyed entity or an aggregation of account holders?
KYC and profile	KYC system, onboarding	Client id	KYC alerts, rubric stated inputs, PEP	Is the risk profile stored as a score, a category, or questionnaire answers?
RM assignment and hierarchy	CRM, HR	Client id, RM id, team id	Row-level security, team page, audit actors	Is assignment history kept, and are cover and checker roles represented?
Accounts and balances	Core banking	Account id, client id	Cash class, liquidity coverage	Cleared versus ledger balance; intraday visibility?
Custody positions	Securities or custody system	Portfolio id, instrument id	Snapshots, allocation, concentration, look-through, lending values	Is a point-in-time snapshot kept daily, or only current state?
Orders and transactions	Order management, execution, settlement	Transaction id, portfolio id	Attribution flows, behavioural features, decision debt	Are orders (placed, amended, cancelled) available, or only settled trades?
Instrument master and prices	Security master, market data	ISIN or internal id	Valuation dates, sector and region, liquidity tier, stale marks	Which vendor, which close, which cut-off per market?
Structured product terms	Structuring or product system	Instrument id	Look-through legs, barriers, autocall dates	Are term sheets structured data or documents?
Funds and NAVs	Transfer agent, fund houses	Fund id	Liquidity tiers, stale valuation, commitments	What is the NAV publication lag per fund type?
Lending and collateral	Lending or collateral system	Facility id, portfolio id	LTV headroom, margin-call alerts, collateral actions	When is LTV recomputed and where are waivers recorded?

DOMAIN	TYPICAL SYSTEM	KEY	WORKBENCH USE	ASK
Mandates	Portfolio management system	Mandate code, portfolio id	Bands, exclusions, breach alerts	Are client-directed deviations flagged in data?
Corporate actions	Custodian, CA system	Event id, instrument id	Dated decisions in the call plan	Are election deadlines and client elections captured?
Notes, calls, messages	CRM, recorded lines, secure mail, portal	Client id, timestamp	Call plan cadence, promises, sensitivities, unanswered-contact alert	Are notes structured (channel, date) and in which languages?
Cash needs and life events	CRM fields, financial planning tool	Client id	Liquidity coverage, moments ahead, call plan deadlines	Are needs dated and amount-bearing, or prose?
Research and market events	Research platform, market data	Event id	Signals, impact, idea desk, morning brief	Is there a feed, or is it documents and email?
Product approvals and lists	Product catalogue, approved list	Instrument id	Idea universe, suitability gate	Which list may the RM propose from?

S-1 **Team & Data Platform** For each domain, which system is the system of record and which hold copies? When they disagree, which wins?

Why we ask Every figure on screen links to its rows; the rows must come from one authoritative place.

S-2 **Team & Data Platform** Are identifiers stable across systems and over time, and is there a mapping service for client, account, portfolio and instrument ids?

Why we ask Every derived table is keyed by client id; a broken join silently drops a household from the book.

S-3 **Team & Data Platform** Which fields are personal or restricted, and what may leave the bank to a model provider? Is pseudonymisation acceptable?

Why we ask Governs what the model may see when writing briefs and drafts, and what the traces may store.

S-4 **Team & Data Platform** Who owns each system, who can grant read access, and who would own a new event feed from it?

Why we ask Every integration in section 8 needs a producer-side owner; without one, files are the realistic path.

Recency, quality and coverage per system

This is where the six o'clock question is answered. For every source: when is it final, when can we read it, how often is it wrong, and what share of the book does it cover.

SOURCE	RECENCY: ASK	QUALITY: ASK	COVERAGE: ASK	EFFECT ON THE 6 A.M. BRIEF
Core banking balances	Real-time posting; when is end-of-day final and readable?	Reversals and re-postings frequency	All currencies and booking centres?	Liquidity coverage can be current if readable overnight
Custody positions	Position close time per market; file or query availability time	Failed settlements, pending trades, corrections	All custodians and sub-custodians?	The single largest blocker: if positions are 9 a.m., the brief works from the previous close and refreshes
Prices and FX	Close time per market; late prices; FX cut-off	Stale, missing and vendor-flagged prices	All instrument types incl. private and structured?	Valuation dates per line decide which figures the brief may quote
Fund NAVs	Publication lag per fund type	Restatements	All fund houses?	Alternatives stay at last NAV with date shown
Orders and transactions	Real-time, intraday feed or end-of-day extract?	Amend and cancel handling	All channels: RM, portal, mobile?	Yesterday's activity is in the brief only if the feed is intraday or early
Lending and LTV	Recomputation time; dependency on prices	Manual overrides, waivers	All facility types?	Margin-call proximity is only as current as LTV
Client profile and KYC	On change; nightly extract?	Missing fields, expired reviews	All segments?	KYC alerts and cross-border rules need residence and dates to be current
RM assignments	On change; how fast must access follow?	Orphan clients, duplicate primaries	Every client assigned?	Isolation is only as good as this feed
Notes and messages	As written; timestamps reliable?	Empty or back-dated notes	Which channels are captured?	Call plan cadence, promises and unanswered messages
Corporate actions	Announcement to availability lag	Amended events	All markets?	Election deadlines on the call sheet
Research and events	Publication time; overnight events before 6 a.m.?	Corrections	Which asset classes and regions?	Signals in the brief need events published before the run

R-1

Team & Data Platform

For each source, at what time is yesterday's data final, at what time can a consumer read it, and in which time zone? Where is the gap between the two, and why?

Why we ask The difference between "final at 02:00" and "readable at 09:00" is usually a scheduling choice, not a system limit; it is the three hours we need.

R-2

Team & Data Platform

Can any source deliver incrementally (changes only) instead of a full file, and earlier than the consolidated batch?

Why we ask Incremental deliveries let the workbench apply overnight changes and be ready at six, then reconcile when the full batch lands.

R-3

Team & Data Platform

Which values can be restated after publication (late prices, corrected trades, NAV restatements) and how is a restatement signalled?

Why we ask A brief written at six from figures restated at nine must show that it was refreshed, and the audit trail must keep both.

R-4

Team & Data Platform

What data-quality checks already run per source, where are failures visible, and who fixes them?

Why we ask The workbench's thirteen checks and dead-letter view should extend the bank's controls, and someone must own the failures.

R-5

Team & Data Platform

What share of clients, portfolios and instruments does each source cover? Which clients or products are outside it (external custody, held-away assets, insurance)?

Why we ask A household total that omits held-away assets misleads the RM; gaps should be shown as gaps.

R-6

Team & Data Platform

How much history is online per source and are daily point-in-time snapshots kept?

Why we ask Momentum, stress windows, behavioural features and the replay clock need dated history, not just current state.

R-7

Team & Data Platform

Is the client's reporting-currency conversion done at source or expected downstream, with which rate and time?

Why we ask Household totals and FX effects are only comparable under one convention.

Event or micro-batch support

The workbench applies typed messages idempotently from Kafka or from files. Whether a source can emit events, deliver micro-batches, or only a daily file decides how early the brief can be ready and how current it stays through the day.

V-1

Team & Data Platform

Does each source publish business events today (message bus, webhooks, outbox), expose change data capture, deliver micro-batches at intervals, or only a daily file?

Why we ask Decides whether the Kafka front door or the batch adapter is the first integration per source, and what "current" can mean.

V-2

Team & Data Platform

For sources that emit events: is there a stable idempotency key and a business event time separate from publish time?

Why we ask The landing zone dedupes on the key; the change log orders on event time; replay must be safe.

V-3

Team & Data Platform

Are events ordered per client or account, and can a consumer partition by client id?

Why we ask The consumer partitions by client so a household's changes apply in order.

V-4

Team & Data Platform

How are late, out-of-order and corrected events represented?

Why we ask Restatements must be applied, not appended; the brief must say it was refreshed.

V-5

Team & Data Platform

Can topics be replayed from an offset or date, how long is retention, and is there a bootstrap path (snapshot topic or initial file)?

Why we ask A new consumer, a new RM book or a rebuilt vector store needs a cold start.

V-6

Team & Data Platform

For micro-batch sources: at what intervals, and can the interval be shortened overnight so the last batch lands before 05:30?

Why we ask Micro-batches every hour overnight would let the six o'clock run see nearly everything without any real-time integration.

V-7

Team & Data Platform

Which events carry personal data, and is tokenisation applied on the bus?

Why we ask Determines what may sit in the landing zone and what must be masked before a model sees it.

V-8

Team & Data Platform

Who owns topic naming, schema registry and versioning, and who triages dead letters: producer or consumer team?

Why we ask Message contracts are versioned; rejected messages appear in a dead-letter view that someone must own.

V-9 **Team & Data Platform** What volumes and peaks per source (end of day, month end, corporate-action seasons)?

Why we ask Sizes consumer batches and decides between book-wide and per-client rebuilds of vectors and alerts.

V-10 **Team & Data Platform** Is there a production-like test environment with masked replay or synthetic feeds?

Why we ask Ingestion must be proven against real shapes before the first six o'clock run.

ASSET CLASS	EVENTS TO ASK FOR	IF ONLY FILES EXIST, DETECT BY	WHAT THE RM SEES
Cash and FX	Balance posted; deposit maturing; accumulator fixing; large payment	Nightly balance diff; maturity calendar	Liquidity coverage; events as moments
Equities and IPO	Order placed, executed, cancelled; settlement; RSP executed or missed; IPO allotted; corporate action announced with deadline	End-of-day extract; CA calendar	Concentration alerts; debt; dated election call sheet
Funds and alternatives	NAV published; capital call; distribution; gate change	Transfer-agent file; commitment schedule	Stale-valuation alert; dated needs
Bonds and structured products	Coupon; call announced; maturity; rating change; barrier or autocall event	Master and term-sheet calendar; price diff	Income; signals on household; look-through changes
Lending	LTV recomputed; margin call issued, cured, waived; drawdown	End-of-day facility snapshot	Margin-call proximity; projected breach date
Client and relationship	Onboarded; profile or KYC changed; RM assignment changed; note written; message received; cash need added	Nightly CRM extract	Isolation updated; promises; unanswered messages
Market	Event published; house view changed; threshold move	Daily research digest	Signals, impact, the market sentence

Agent build: content, ideas and communication

Imagine the RM at six. The brief is written, the calls are ranked, each client's story is ready. What else would she want the machine to have done overnight? Written the house view into words for each client. Turned this week's idea into the six clients it fits, each with a draft in their language. Prepared the messages she will send at nine, through the channel each client prefers, already passed through compliance. These questions establish whether the bank has the platforms and permissions for that, and what the RM would actually do with it.

Content authoring and publishing

G-1 **Ranked Actions & Ideas** Does the bank have a content authoring and publishing platform for research and Chief Investment Office insights? Who writes, who approves, in which formats and languages, and is there a feed the workbench can read?

Why we ask House view is the source of signals and ideas; if it exists as documents and email only, the workbench must extract it, and approval status must travel with it.

G-2 **Workflow, Outreach & Governance** Is there a demand for personalised content per client: the house view rewritten for their holdings, their language and their profile, with the parts that do not apply removed?

Why we ask The journey already writes prose from a client's facts; extending it to CIO content is a small step if compliance accepts model-rewritten text under approval.

G-3 **Workflow, Outreach & Governance** Which content may an RM adapt and which must be sent verbatim? Who signs off a personalised version, and how is the derivation from the approved original evidenced?

Why we ask The gateway appends disclosures and logs edits today; personalised content needs a lineage from the approved source.

G-4 **Ranked Actions & Ideas** Are trade ideas produced centrally (advisory desk, product specialists) or by RMs? How are they distributed, how is client fit assessed, and how is uptake measured?

Why we ask The idea desk inverts ideas to clients with a suitability gate and tracks drafted and sent per idea; central ideas should enter it as data, not as a PDF.

G-5 **Ranked Actions & Ideas** What does the RM do today with a new idea: read it, pick clients from memory, write the message by hand? How long does that take per idea?

Why we ask That is the manual step the idea desk removes; the time saved is the business case.

Communicating with clients

G-6 **Workflow, Outreach & Governance** Through which channels may the bank contact clients: email, secure portal, mobile app notification, WhatsApp or other messaging, SMS, phone, letter? Which are approved for advice, which for information only, and per which jurisdiction?

Why we ask The gateway checks every message; the channel decides which checks and disclosures apply and whether an agent may ever send.

G-7 **Workflow, Outreach & Governance** Where is client channel preference and consent recorded, including opt-outs and quiet hours?

Why we ask Drafts should default to the client's preferred channel and language; consent gates any agent-sent message.

G-8 **Workflow, Outreach & Governance** Is WhatsApp or another messaging platform in use with clients today, officially or unofficially, and is it archived and surveilled?

Why we ask If it is the real channel, the workbench should draft for it and log it; if it is unofficial, that is a conduct finding for the team page, not a feature.

G-9 **Workflow, Outreach & Governance** Does the mobile app support personalised content or notifications per client, and can a message drafted by the RM be published into the client's app inbox?

Why we ask An app inbox is the cleanest channel for agent-sent informational content, with read receipts and no forwarding risk.

G-10 **Workflow, Outreach & Governance** Which templates and disclosures must a client message carry by product, residence and channel, and who approves templates?

Why we ask The gateway appends required disclosures; the bank's template library should replace the reference file.

G-11 **Workflow, Outreach & Governance** Is there an appetite for an agent to send low-risk informational messages itself (KYC reminder, meeting scheduling, market update, "since we last spoke" summary) with the RM notified, and what evidence would compliance want first?

Why we ask Shadow mode measures agreement between what an agent would say and what the RM said; the bank sets the threshold that earns that step.

G-12 **Workflow, Outreach & Governance** What must be logged when a message is sent: content, channel, time, approver, disclosures, evidence used? Where must it land: the workbench, the CRM, the archive?

Why we ask Every send is an audit event today; the bank may need it mirrored into its own archive for surveillance.

What the RM would do with the agents

G-13 **Daily Planning & Conversations** If the RM opened the workbench at six and found the brief, the call sheet, six personalised drafts and three proposed tasks, what would she change before acting, and what would she want to be asked first?

Why we ask Defines the confirm points: which outputs are ready to use, which need a glance, which need a decision.

G-14 **Daily Planning & Conversations** Which questions does the RM ask her assistant or her systems most often, in her words?

Why we ask The assistant's grammar is a list of RM questions; the bank's list makes it useful on day one without a model.

G-15 **Daily Planning & Conversations** During a client call, what does the RM wish she could look up in seconds but cannot today?

Why we ask The companion answers cues in a second from the engines; the wish list tells us which engines matter and which are missing.

G-16 **Workflow, Outreach & Governance** Which promises does the RM make most often, and how does she track them today?

Why we ask The promise ledger extracts commitments from notes and drafts their fulfilment; the common ones deserve templates.

G-17 **Client Intelligence & 360** What would make the RM trust a machine-written client story enough to read it aloud on a call?

Why we ask Every sentence links to its rows; if trust needs something else (a human check, a house-approved phrase), the design should know.

G-18 **Team & Data Platform** What would the team head want to see before allowing an agent to talk to a client at all, and what single switch would she want in her hand?

Why we ask The machine panel shows agreement per playbook and a pause-all switch; the head's answer sets the threshold and the evidence.

Calls to action and tasks

Information tells the RM what is true; insight tells her what it means; a call to action tells her what to do; a task is that action with an owner, a due date and evidence of completion. The workbench turns every insight into a proposed task the RM confirms, and every confirmation into an audit event. These questions establish what the bank's RMs actually do with what they learn, where tasks live today, and which ones a machine may prepare, propose or complete.

From insight to action: what the workbench proposes today

INSIGHT	CALL TO ACTION	TASK THE RM CONFIRMS	EVIDENCE
LTV headroom falling toward the trigger	Call the client before the breach	Discuss collateral: top-up, reduce drawing or de-risk; second RM check	Approval and decision log; logged; top facility snapshot
Concentration above limit (direct or look-through)	Propose a trim or a client-directed acknowledgement	Trade idea with suitability gate; or record the client's instruction to hold	Idea approved with note; client flag recorded
Cash need approaching with thin coverage	Agree the funding route before the date	Liquidity plan: sell, draw, or reschedule; lending conversation	Plan noted; transaction change observed
Severe market signal reaching the household	Call the exposed clients first	Discuss the impact; send a market update in the client's language	Call logged; through the gateway
Stated versus observed risk mismatch	Re-profile or document the difference	Rubric override with reason, or a review meeting	Override with reason logged; local file attached to file
KYC review due or overdue	Schedule the refresh	KYC refresh playbook; documents requested	Meeting scheduled; system updated
Client message unanswered	Reply today	Draft reply in the client's language; send through the gateway	Sent and logged

INSIGHT	CALL TO ACTION	TASK THE RM CONFIRMS	EVIDENCE
Promise made on a call	Fulfil it by the date	Send the material; the draft is attached to the promise	Promise made; the send event
Decision agreed but never executed	Bring it back to the client	Talking point on the next call; staged deployment proposal	Execution confirmed; client's decision
Contact overdue against cadence	Call, even with nothing urgent	Relationship call; review scheduling	Call logged; review
Idea fits several clients	Reach them, in order	One draft per client; reverse-enquiry record where required	Drafts sent; per idea

Questions about actions and tasks

T-1 **Ranked Actions & Ideas** For each insight the workbench produces, what does an RM do about it today, and what would the bank want her to do? Are there insights that should produce no action at all?

Why we ask The ranked actions and the call plan encode a mapping from insight to action; the bank's own expectation replaces our defaults, and some alerts may be for information only.

T-2 **Workflow, Outreach & Governance** Where do RM tasks live today: a CRM task list, a case-management tool, email, a spreadsheet, the RM's head? Who assigns them, and who else can see them?

Why we ask The workbench proposes tasks and logs their completion; if the bank has a task system, tasks must be written there too, or the RM will keep two lists.

T-3 **Workflow, Outreach & Governance** Which tasks come to the RM from elsewhere: compliance (KYC, documentation), operations (failed settlement, missing form), the team head, the client, product desks? In what form and with what deadline?

Why we ask The six o'clock task list should include what others need from the RM, not only what the data suggests; these need a feed or an inbox.

T-4 **Daily Planning & Conversations** How are tasks prioritised today when there are more than the day allows, and what happens to the ones that slip?

Why we ask The call plan packs by deadline and priority and shows overflow explicitly; the bank's rules for what may slip decide the packing order.

T-5 **Workflow, Outreach & Governance** What counts as a task being done, and what evidence must exist: a note, a system update, a sent message, an approval?

Why we ask The workbench marks tasks done from the event that proves it (send, decision, log) rather than a tick; each task type needs a defined proof.

T-6 **Daily Planning & Conversations** Which tasks have a regulatory or service-level deadline (KYC, margin call notice, complaint response, corporate-action election), and where are those deadlines held?

Why we ask Dated deadlines drive the clock term of the call plan and the "returned early" rule on deferred clients.

T-7 **Workflow, Outreach & Governance** Which tasks may a machine complete outright (assemble documents, pre-fill a form, schedule from a calendar, draft a message), which may it propose for one-click confirmation, and which must the RM originate?

Why we ask This is the autonomy line per task type: complete, propose, or prepare only. The workbench proposes today and confirms per card.

T-8 **Workflow, Outreach & Governance** May a task be assigned to someone other than the RM (assistant, operations, specialist), and may the machine route it there directly?

Why we ask Some calls to action are handoffs; routing needs the team model and the recipient's task system.

T-9 **Team & Data Platform** How does the team head see task load and slippage across the team today, and what would she want escalated to her automatically?

Why we ask The team page shows calls today, deferrals and approvals waiting; escalation rules decide what reaches her without asking.

T-10 **Daily Planning & Conversations** When an RM defers or dismisses a proposed task, must the reason be recorded, and does anyone review the pattern?

Why we ask Deferrals with reasons are logged and feed both the call plan and the conduct panel; the bank decides whether the pattern is reviewed.

T-11 **Client Intelligence & 360** Which task outcomes should feed back into client data: a client's stated intention, a waiver, a preference, a promise kept?

Why we ask Closing a task often creates a fact (client chose to hold, prefers WhatsApp, will fund by Friday); those facts should land in structured fields, not only in a note.

T-12

Daily Planning & Conversations

What would the RM want the six o'clock task list to look like: how many items, in what order, with what attached, and what should never be on it?

Why we ask Defines the shape of Today: the brief, the call sheet and the tasks are one page, and it must be readable in ninety seconds.

The fifteen to ask first

If the first workshop has ninety minutes, these decide whether the six o'clock brief is possible and what it may contain.

- 1 For custody positions and prices: at what time is yesterday final, at what time can we read it, and why the gap? (R-1)
- 2 Which sources can deliver incrementally or in overnight micro-batches rather than one 9 a.m. file? (R-2, V-1, V-6)
- 3 Which system is the golden source for the client and household, and for RM-to-client assignment with history? (S-1, E-1)
- 4 Which identity provider issues RM identities, and which claims carry role and team? (E-4)
- 5 Who else may see a client besides the primary RM, and is a team head's view the whole team or aggregates? (E-2, E-3)
- 6 Which portfolios are discretionary, advisory, execution-only or custody, and where is that flag? (A-E1)
- 7 Where are mandate bands, limits and exclusions codified, and do concentration limits look through? (A-L3, A-E2)
- 8 How often is LTV recomputed, and where are triggers and waivers recorded? (A-L2)
- 9 Are structured notes decomposed into legs, and are lending values carried per holding? (A-B1, A-B3)
- 10 Where do notes, calls, client messages and dated cash needs live, in which languages? (Section 6, R-5)
- 11 What is the bank's risk-profile scale, and may stated be compared with observed? (A-L4)
- 12 Which suitability, appropriateness and cross-border rules must sit in front of every idea and message, and in what form can compliance supply them? (A-B4, A-X1)
- 13 Is there a research and CIO content platform with a feed and approval status, and is personalised content wanted? (G-1, G-2)
- 14 Which client channels are approved, where is consent recorded, and is an agent-sent informational message conceivable? (G-6, G-7, G-11)

15

Where do RM tasks live today, what proves a task done, and which tasks may a machine complete rather than propose? (T-2, T-5, T-7)

The data model behind the RM

The model the workbench runs on today, with the identifiers the bank must supply or agree. Four schemas: raw holds what the bank's systems supply, derived holds what the engines compute and what people decide, access holds identities, ingest holds the landing zone. Every row a person decides carries the Client ID, so row-level security and the audit trail follow the client wherever the row lives. Household is the one entity we propose adding: today the client is the household.

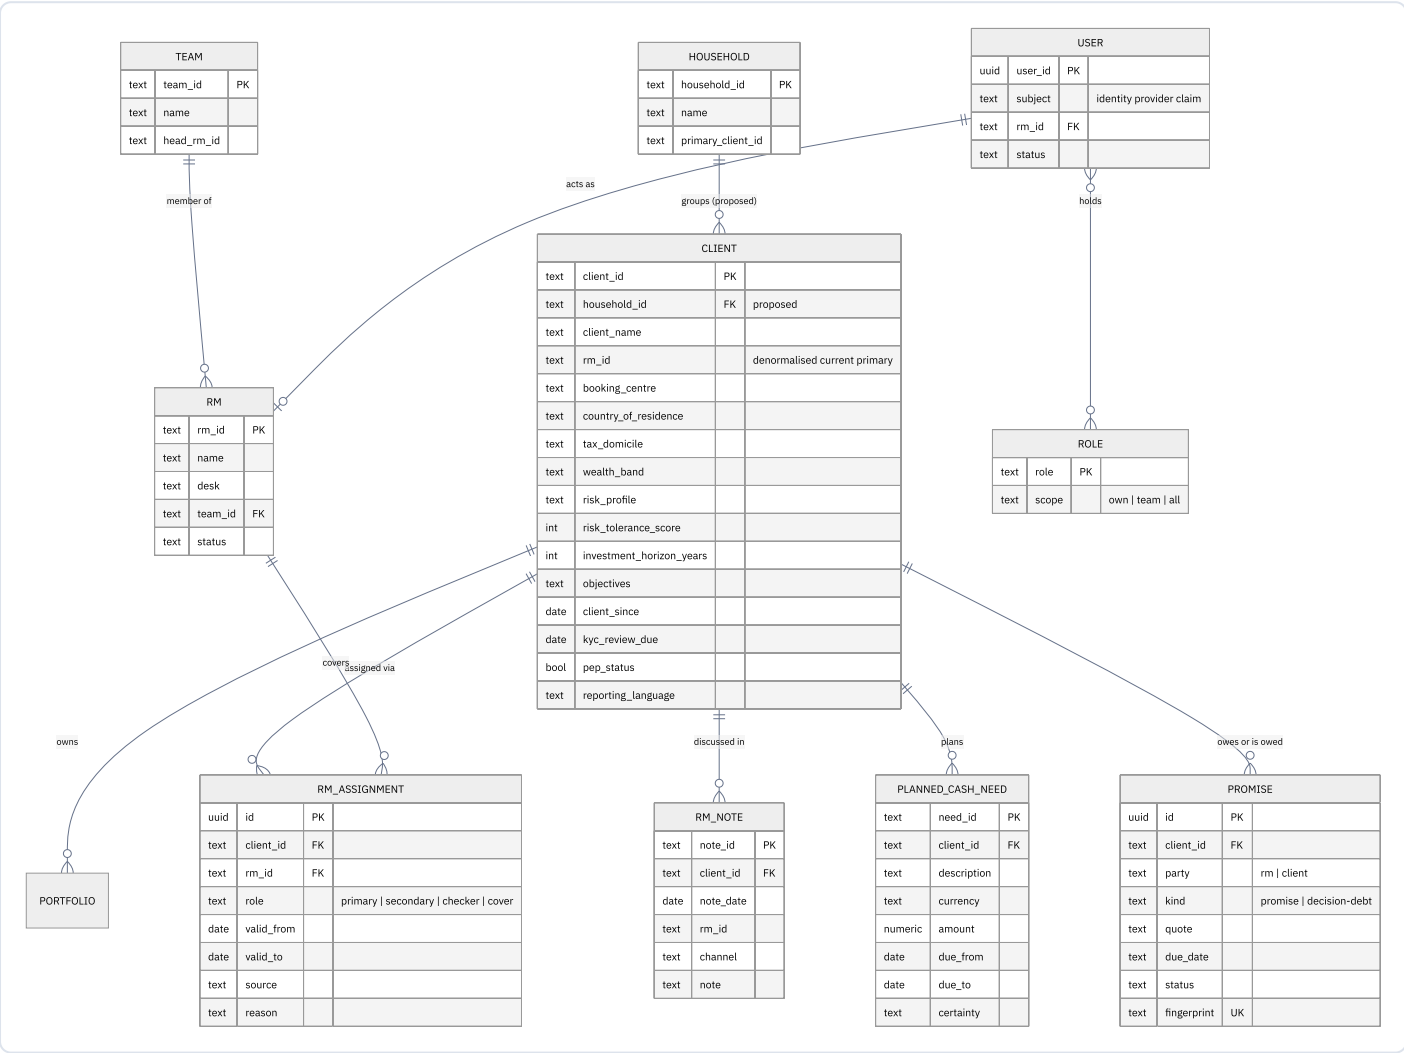
Identifiers

IDENTIFIER	MEANING	OWNER SYSTEM TO CONFIRM	USED TO JOIN	ASK
Household ID	The economic unit the RM manages: joint holders, family entities, trusts under one relationship	CRM or client master (proposed)	Household → Clients	Does the client have this or is the client the household?
Client ID	The legal person or entity with a profile, KYC and reporting language; the key every derived row carries	Client master	Everything client-owned	Is it stable across systems over time?
Portfolio ID	A mandate-governed container of holdings with a service model and base currency	Portfolio management or custody	Client → Portfolios → Holdings, Transactions, AUM	Is a portfolio one cust account or several?
Account ID	A cash or custody account in core banking	Core banking	Portfolio → Accounts (proposed link)	How do accounts link to portfolio and client?
Instrument ID	The security master key (ISIN or internal), carrying asset class, sector, region, liquidity tier	Security master	Holdings, Prices, Transactions, Look-through legs	ISIN, internal id, or both? how are structured notes identified?

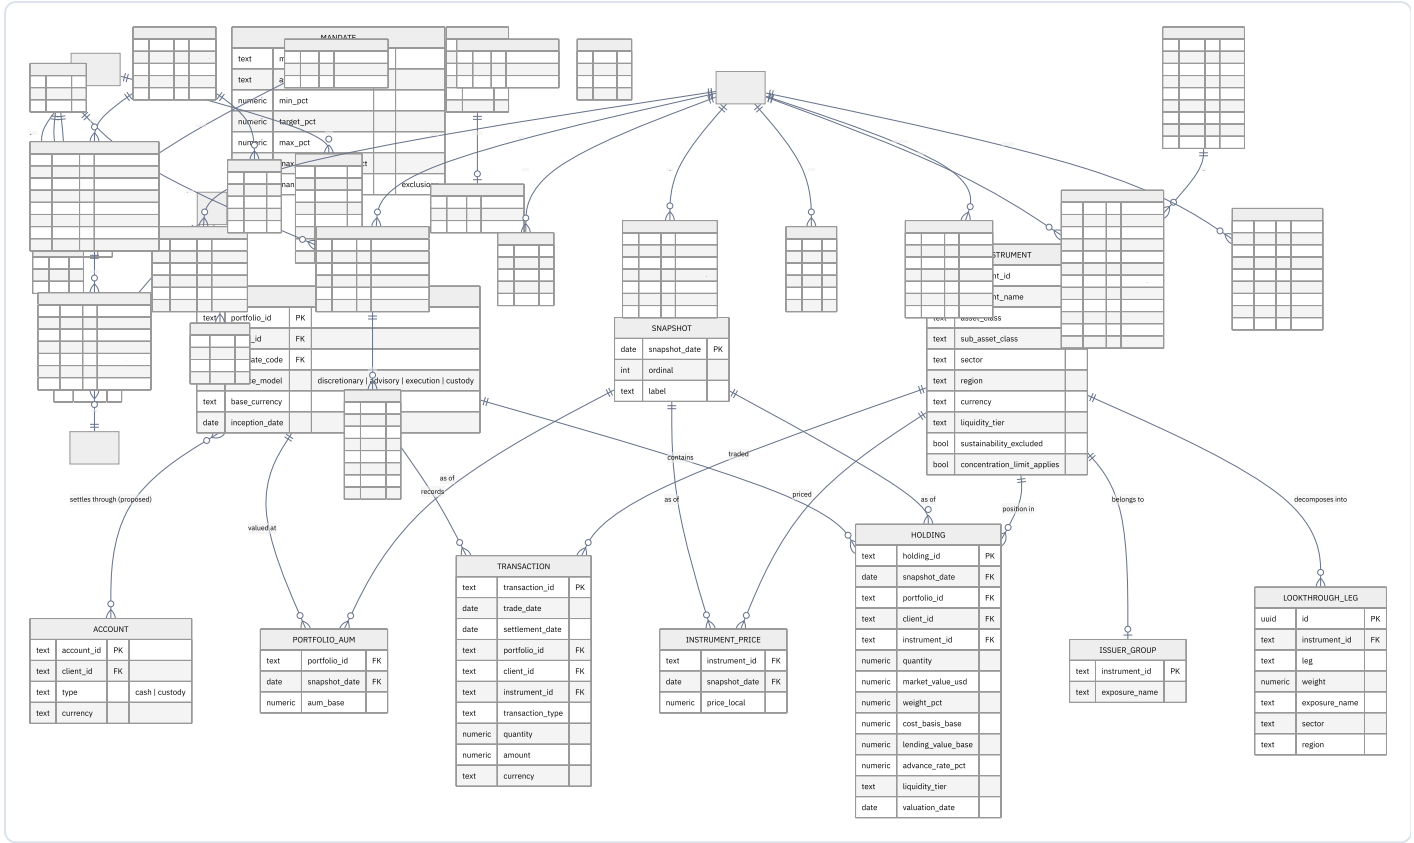
IDENTIFIER	MEANING	OWNER SYSTEM TO CONFIRM	USED TO JOIN	ASK
Snapshot date	The valuation date positions, prices, AUM and LTV are stated at; the time axis of the whole model	Custody close	Holdings, Prices, Portfolio AUM, Facility snapshots, Market context	Daily or more frequent point-in-time snapshots kept?
Transaction ID	One order, execution, income or fee event	Order management or settlement	Portfolio, Client, Instrument	Are amendments new ids or updates?
Facility ID	A credit facility with a collateral portfolio and trigger levels	Lending system	Client, Collateral portfolio, Facility snapshots	One facility per portfolio or many-to-many?
Mandate code	The strategic allocation with bands per asset class and exclusions	Portfolio management	Portfolios	Are client specific mandates versioned?
RM ID • Team ID	The relationship manager and the team or desk	HR or CRM	Assignments, Users, Notes	Which system is authoritative for reassignment?
Note ID • Need ID • Commitment ID	An interaction record; a dated cash need; an uncalled commitment	CRM, planning tool, alternatives administrator	Client	Are needs structured with data amount?
Event ID • Series ID	A market event; a market data series	Research, market data	Signal rules, Impact runs	Is there a curated market data feed?

IDENTIFIER	MEANING	OWNER SYSTEM TO CONFIRM	USED TO JOIN	ASK
Run IDs • Trace IDs	Load runs, vector runs, impact runs, model traces	Workbench	Every computed row back to its inputs	None; th are ours.

Client, household and relationship



Portfolios, positions and transactions



Lending, commitments and market context

What the engines compute and what people decide

Ingestion, reference and policy

Row-level security applies to every table that carries a Client ID in raw and derived; access and ingest metadata are exempt because the entitlementment function reads them. Assignments, users and roles are the access model; the RM columns on the client are a denormalised copy of the current primary, kept in step by the loader.

Discovery questionnaire for the agentic RM workbench in Private Banking and Retail Wealth. Category tags: C1 Client Intelligence & 360, C2 Market Signals & Risk, C3 Ranked Actions & Ideas, C4 Workflow, Outreach & Governance, C5 Daily Planning & Conversations, C6 Team & Data Platform Foundation. Question ids: C catalogue to confirm, J what the RM does, E entitlements, A asset classes, I data and insights, S source systems, R recency and quality, V events and micro-batches, G agent build, T calls to action and tasks. The workbench modules referenced are those built in iterations 0 to 10 and the data platform (ADR-001 to ADR-016). Prepared 18 Sep 2026.